

Exam CS4400: Deep Reinforcement Learning

practice early! | anytime

Student name: _____

Student number: _____

- This is a closed-book individual examination with **9 questions** and a total of **42 points**.
- Do not open the exam before the official start of the examination.
- If you feel sick or otherwise unable to take the examination, please indicate this *before* the exam starts.
- The examination lasts **150 minutes** after the official start.
- This gives you roughly 3 minutes per point. Use your time carefully!
- You can hand in your exam solution any time until 15 minutes before the end of the exam and leave the examination room quietly. In the last 15 minutes, no one can leave the examination room to help other students concentrate on finishing their exam.
- Only one student can visit the bathroom at the same time. In the last 15 minutes, no bathroom visits are possible.
- Use of course books, readers, notes, and slides is **not** permitted
- Use of (graphical) calculators or mobile computing devices (including mobile phones) is **not** permitted.
- Write down your name and student number above.
- Write your **student number on each sheet** of the exam after the exam started.
- You can write your answer on the free space under each question.
- If you need more space, use the back of another exam-sheet and write where to find the answer under the question. Ask for additional empty pages if you need them.
- Use pens with black or blue ink. Pencils and red ink are not allowed!
- Clearly cross out invalid answers. If two answers are given, we consider the one with less points!
- Write clearly, use correct English, and avoid verbose explanations. Giving irrelevant information may lead to a reduction in your score.
- This exam covers all information on the slides of the course, the tutorials and everything discussed in lectures.
- This exam assumes a familiarity with the stated background knowledge of the course.
- The total number of pages of this exam is 5 (excluding this front page).
- Exam prepared by Wendelin Böhmer. ©2022 TU Delft.

Question 1 (multiple choice):**(10 points)**

Please mark only the correct answers with a **cross** like this: ☒. If you wish to **unmark** a marked answer, **fill** the entire square and **draw an empty** one next to it like this: ☐ ■

Only one answer per question is correct. You will receive 1 point per correct answer, except if multiple squares are marked. Wrong answers yield no points, but are also not punished. Good luck!

1.1: Recurrent neural networks *without* gates can **not** have:

- ☐ exploding gradients
- ☐ vanishing gradients
- ☐ memory cells
- ☐ non-linear transfer-functions

1.2: Which of the following is **not** a normalization layer module in PyTorch?

- ☐ `torch.nn.LayerNorm`
- ☐ `torch.nn.TemporalNorm`
- ☐ `torch.nn.InstanceNorm2d`
- ☐ `torch.nn.BatchNorm1d`

1.3: We use double Q-learning to counter which problem for bootstrapping?

- ☐ $\mathbb{E}[\max_a Q(s, a)] \leq \max_a \mathbb{E}[Q(s, a)]$
- ☐ $\mathbb{E}[\max_a Q(s, a)] \geq \max_a \mathbb{E}[Q(s, a)]$
- ☐ $\mathbb{E}[Q(s, \arg \max_a Q'(s, a))] \leq \max_a \mathbb{E}[Q(s, a)]$
- ☐ $\mathbb{E}[Q(s, \arg \max_a Q'(s, a))] \geq \max_a \mathbb{E}[Q(s, a)]$

1.4: In which of the following cases is a *dueling network architecture* useful?

- ☐ When the value of two actions are too similar.
- ☐ When maximization introduces overestimation bias.
- ☐ When a pessimistic value estimate is needed.
- ☐ When many actions have the same successor state.

1.5: What is the correct normalization factor x for $Q(\lambda)$ targets $Q_\lambda^*(s_t, a_t) := x \sum_{n=0}^m \lambda^n Q_{n+1}^*(s_t, a_t)$?

- ☐ $x = 1 - \lambda^{m+1}$
- ☐ $x = \frac{1}{1 - \lambda^{m+1}}$
- ☐ $x = \frac{1 - \lambda^{m+1}}{1 - \lambda}$
- ☐ $x = \frac{1 - \lambda}{1 - \lambda^{m+1}}$

1.6: Which uncertainty measure can **not** be used as an estimate of the value function's variance?

- ☐ random network distillation
- ☐ pseudo-counts
- ☐ random hash functions
- ☐ ensembles

1.7: Which of the following statements about deep exploration is correct?

- ☐ deep exploration converges faster to the optimal policy than ϵ -greedy
- ☐ deep exploration requires to propagate future uncertainty
- ☐ deep exploration requires to count how often a specific state has been seen
- ☐ deep exploration requires deep neural networks

1.8: Which value learning method is the most *stable*?

- ☐ independent learning
- ☐ on-policy learning
- ☐ off-policy learning
- ☐ offline learning

1.9: Which algorithm uses methods from offline RL to stabilize learning?

- ☐ DQN
- ☐ REINFORCE
- ☐ TRPO
- ☐ SAC

1.10: Which MARL algorithm does **not** use centralized training?

- ☐ IQL
- ☐ MADDPG
- ☐ COMA
- ☐ QMIX

Question 2:

(2 points)

Describe in 4 sentences or less **two** examples where a value function helps a policy gradient algorithm. The examples must come from *different* algorithms.

Question 3:

(2 points)

In 4 sentences or less, name and explain **two** examples where *robust reinforcement learning* is useful.

Question 4:**(2 points)**

Explain in 4 sentences or less the difference between *Thompson sampling* and *optimistic exploration*.

Question 5:**(3 points)**

Give the joint actions of all Nash-equilibria for a two-player single-state general-sum game with the following reward matrix, where entry x/y denotes the reward x for player 1 and y for player 2:

P1/P2	a_1^2	a_2^2	a_3^2	a_4^2
a_1^1	-2/1	4/-2	2/3	0/-1
a_2^1	-1/5	3/1	0/4	-3/2
a_3^1	1/2	2/-2	-1/1	2/1
a_4^1	0/2	-3/3	1/4	5/1

Which Nash equilibrium would player 1 prefer? Which would be better for player 2?

Question 6:**(4 points)**

Let $v := \sum_{t=0}^{\infty} \gamma^t r_t$ denote the value in a MDP with a single state and action, where the reward $r_t \sim \mathcal{N}(\mu, \sigma^2)$ is drawn i.i.d. from a normal distribution and $\gamma \in (0, 1)$ denotes the discount factor. *Without* using the fact the variance of a sum of independent variables is the sum of the variables' variances, prove analytically that the variance of v is

$$\mathbb{V}[v] = \frac{\sigma^2}{1 - \gamma^2}.$$

Question 7:**(5 points)**

Derive the vanilla actor-critic policy gradient $\nabla_{\theta} \mathcal{L}_{\pi}[\theta]$ of policy π_{θ} from the off-policy gradient $\nabla_{\theta} \mathcal{L}_{\mu}[\theta]$:

$$\nabla_{\theta} \mathcal{L}_{\mu}[\theta] := -\nabla_{\theta} \mathbb{E}_{\mu} \left[\sum_{t=0}^{n-1} \gamma^t A_t^{\pi} \frac{\pi_{\theta}(a_t|s_t)}{\mu(a_t|s_t)} \right] = \underbrace{\dots}_{\text{derivation}} \approx -\mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^{n-1} \gamma^t A_t^{\pi} \nabla_{\theta} \ln \pi_{\theta}(a_t|s_t) \right] =: \nabla_{\theta} \mathcal{L}_{\pi}[\theta],$$

where $A_t^{\pi} := r_t + \gamma V^{\pi}(s_{t+1}) - V^{\pi}(s_t)$ is the advantage at time t , based on the state-value function $V^{\pi}(s)$ that does not depend on θ . Which approximation do you need to make? Make sure every step can be followed easily!

Question 8: (programming)**(6 points)**

You only have to insert the missing code segment at the line(s) marked with `#YOUR CODE HERE`. Please use correct Python/PyTorch code. Singleton dimensions of tensors can be ignored, i.e., you do not need to `(un)squeeze` tensors. If you forget a specific command, you can define it first, both the

signature (input/output parameters) and a short description what it does. Using your own definitions of existing PyTorch functions will not yield point deductions. If no similar PyTorch function exists, your definition will be considered as wrong code and you will not receive the corresponding points.

Implement the following DDPG policy objective efficiently in the given `MyLearner` class:

$$\max_{\theta} \mathbb{E} \left[\frac{1}{m} \sum_{i=1}^m \frac{1}{n_i} \sum_{t=0}^{n_i-1} \gamma^t Q_{\phi}(s_t^i, \pi_{\theta}(s_t^i)) \mid \tau_{n_i}^i \sim \mathcal{D} \right],$$

where $\tau_{n_i}^i := \{s_t^i, a_t^i\}_{t=0}^{n_i-1}$ is a history of states $s_t^i \in \mathbb{R}^d$ and actions $a_t^i \in \mathbb{R}^b$ up to time $n_i - 1$.

Hint: The Q-value module `value` (computes the Q-values Q_{ϕ} for a minibatch) is given and takes the concatenation of a state- and an action-tensor of equal size (except for the last dimension) as input. The policy module `policy` (computes the policy π_{θ} for a minibatch) is given, takes a state-tensor as input and returns an action tensor of the same size (except for the last dimension).

```

1 import torch
2
3 class MyLearner:
4     def __init__(self, value, policy, gamma=0.99):
5         self.value_model = value
6         self.policy_model = policy
7         self.gamma = gamma
8         self.optimizer = torch.optim.Adam(policy.parameters())
9
10    def train(self, batch):
11        """ Performs one gradient update step on the loss defined above.
12            "batch" is a dictionary of equally sized tensors
13            (except for last dimension):
14                - batch['states'][i, t, :] = s_t^i
15                - batch['actions'][i, t, :] = a_t^i
16                - batch['mask'][i, t] = t < n_i """
17        loss = 0
18        # YOUR CODE HERE
19        self.optimizer.zero_grad()
20        loss.backward()
21        self.optimizer.step()
22        return loss.item()

```

Question 9:

(8 points)

A given linear recurrent neural network takes a time series of n input vectors $\mathbf{x}_t \in \mathbb{R}^d$ and outputs a time-series of n vectors $\mathbf{y}_t \in \mathbb{R}^b$. The RNN computes hidden outputs $\mathbf{h}_t \in \mathbb{R}^q$ without non-linearities:

$$\mathbf{h}_t := \mathbf{W}\mathbf{h}_{t-1} + \mathbf{U}\mathbf{x}_t \in \mathbb{R}^q, \quad \mathbf{y}_t := \mathbf{V}\mathbf{h}_t \in \mathbb{R}^b, \quad 1 \leq t \leq n, \quad \mathbf{h}_0 := \mathbf{0}.$$

- (a) [4 points] Prove by induction that the above update equations are equivalent to the multivariate function $g : \{\mathbb{R}^d\}_{t=1}^n \times \mathbb{R}^{q \times q} \times \mathbb{R}^{q \times d} \times \mathbb{R}^{b \times q} \rightarrow \mathbb{R}^{b \times n}$:

$$g(\{\mathbf{x}_t\}_{t=1}^n, \mathbf{W}, \mathbf{U}, \mathbf{V})_{k,m} := \sum_{t=1}^m (\mathbf{V}\mathbf{W}^{m-t}\mathbf{U}\mathbf{x}_t)_k = (\mathbf{y}_m)_k, \quad 1 \leq k \leq b, \quad 1 \leq m \leq n.$$

Hint: start your proof with showing by induction that $\mathbf{h}_m \stackrel{(ind)}{=} \sum_{t=1}^m \mathbf{W}^{m-t}\mathbf{U}\mathbf{x}_t$.

(b) [4 points] Define the *linear function* $f : \mathbb{R}^{\mathcal{J}} \rightarrow \mathbb{R}^{\mathcal{I}}$ that is equivalent to $g(\{\mathbf{x}_t\}_{t=1}^n, \mathbf{W}, \mathbf{U}, \mathbf{V})$:

$$f(\mathbf{z})_i := \sum_{j \in \mathcal{J}} \Theta_{i,j} z_j, \quad \forall \mathbf{z} \in \mathbb{R}^{\mathcal{J}}, \quad \forall i \in \mathcal{I},$$

by defining the index sets \mathcal{J} and \mathcal{I} , and by constructing the inputs $\mathbf{z} \in \mathbb{R}^{\mathcal{J}}$ from $\{\mathbf{x}_t\}_{t=1}^n$ and the parameter matrix/tensor $\Theta \in \mathbb{R}^{\mathcal{I} \times \mathcal{J}}$ from the g 's parameters \mathbf{U}, \mathbf{V} and \mathbf{W} . Make sure f outputs exactly the same as g !