

CS4400

DEEP REINFORCEMENT LEARNING

Lecture 3: Deep Q-Learning

Wendelin Böhmer

`<j.w.bohmer@tudelft.nl>`



21st of November 2023

Content of this lecture



3.1 Value Approximation

3.2 Stabilization Techniques

3.3 Deep Q-Networks


3.1

Deep Q-Learning Value Approximation


3.1 Residual value gradients



- TD-error regression of data set $\mathcal{D} = \{(s_t, a_t, r_t, s_{t+1}, a_{t+1})\}_{t=1}^n$
 - sampled from one or more trajectories based on π



discrete value function



approximated value function

residual algorithms converge, but not necessarily to the correct values, as s_t and s_{t+1} are not independent (Baird, 1995)

3.1 Residual value gradients



- TD-error regression of data set $\mathcal{D} = \{(s_t, a_t, r_t, s_{t+1}, a_{t+1})\}_{t=1}^n$
 - sampled from one or more trajectories based on π
- Mean-squared error between value $v_\theta(s) \approx V^\pi(s)$ and target

$$\mathcal{L}[\theta] := \mathbb{E}_{\mathcal{D}} \left[\left(\underbrace{r_t + \gamma v_\theta(s_{t+1})}_{\text{target } y_t} - v_\theta(s_t) \right)^2 \right]$$

residual algorithms converge, but not necessarily to the correct values, as s_t and s_{t+1} are not independent (Baird, 1995)

3.1 Residual value gradients



- TD-error regression of data set $\mathcal{D} = \{(s_t, a_t, r_t, s_{t+1}, a_{t+1})\}_{t=1}^n$
 - sampled from one or more trajectories based on π

- Mean-squared error between value $v_\theta(s) \approx V^\pi(s)$ and target

$$\mathcal{L}[\theta] := \mathbb{E}_{\mathcal{D}} \left[\underbrace{\left(r_t + \gamma v_\theta(s_{t+1}) - v_\theta(s_t) \right)}_{\text{TD-error}}^2 \right]$$

- SARSA: MSE between Q-value $q_\theta(s, a) \approx Q^\pi(s, a)$ and target
 - also the name of an on-policy control algorithm

$$\mathcal{L}[\theta] := \mathbb{E}_{\mathcal{D}} \left[\underbrace{\left(r_t + \gamma q_\theta(s_{t+1}, a_{t+1}) - q_\theta(s_t, a_t) \right)}_{\text{target } y_t}^2 \right]$$

residual algorithms converge, but not necessarily to the correct values, as s_t and s_{t+1} are not independent (Baird, 1995)

3.1 Residual gradients break causality



- Let's derive the Q-learning update rule for table $\mathbf{Q} \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{A}|}$
 - after observing transition $s_t, a_t \rightarrow r_t, s_{t+1}$ we minimize

$$\mathcal{L}[\mathbf{Q}] := \left(r_t + \gamma Q_{(s_{t+1}, a^*)} - Q_{(s_t, a_t)} \right)^2, \quad a^* := \arg \max_{a \in \mathcal{A}} Q_{(s_{t+1}, a)}$$

- One gradient-descent step $\mathbf{Q} \leftarrow \mathbf{Q} - \frac{\alpha}{2} \nabla_{\mathbf{Q}} \mathcal{L}[\mathbf{Q}]$

$$Q_{(s, a)} \leftarrow Q_{(s, a)} - \alpha \left(r_t + \gamma Q_{(s_{t+1}, a^*)} - Q_{(s_t, a_t)} \right) \left(\gamma \delta_{(s=s_{t+1}, a=a^*)} - \delta_{(s=s_t, a=a_t)} \right)$$

residual algorithms converge, but not necessarily to the correct values, as s_t and s_{t+1} are not independent (Baird, 1995)

3.1 Residual gradients break causality



- Let's derive the Q-learning update rule for table $\mathbf{Q} \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{A}|}$
 - after observing transition $s_t, a_t \rightarrow r_t, s_{t+1}$ we minimize

$$\mathcal{L}[\mathbf{Q}] := \left(r_t + \gamma Q_{(s_{t+1}, a^*)} - Q_{(s_t, a_t)} \right)^2, \quad a^* := \arg \max_{a \in \mathcal{A}} Q_{(s_{t+1}, a)}$$

- One gradient-descent step $\mathbf{Q} \leftarrow \mathbf{Q} - \frac{\alpha}{2} \nabla_{\mathbf{Q}} \mathcal{L}[\mathbf{Q}]$

$$Q_{(s, a)} \leftarrow Q_{(s, a)} - \alpha \left(r_t + \gamma Q_{(s_{t+1}, a^*)} - Q_{(s_t, a_t)} \right) \left(\gamma \delta_{s=s_{t+1}}^{a=a^*} - \delta_{s=s_t}^{a=a_t} \right)$$

$$Q_{(s_t, a_t)} \leftarrow (1 - \alpha) Q_{(s_t, a_t)} + \alpha \left(r_t + \gamma Q_{(s_{t+1}, a^*)} \right) \quad \checkmark$$

residual algorithms converge, but not necessarily to the correct values, as s_t and s_{t+1} are not independent (Baird, 1995)

3.1 Residual gradients break causality



- Let's derive the Q-learning update rule for table $\mathbf{Q} \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{A}|}$
 - after observing transition $s_t, a_t \rightarrow r_t, s_{t+1}$ we minimize

$$\mathcal{L}[\mathbf{Q}] := \left(r_t + \gamma Q_{(s_{t+1}, a^*)} - Q_{(s_t, a_t)} \right)^2, \quad a^* := \arg \max_{a \in \mathcal{A}} Q_{(s_{t+1}, a)}$$

- One gradient-descent step $\mathbf{Q} \leftarrow \mathbf{Q} - \frac{\alpha}{2} \nabla_{\mathbf{Q}} \mathcal{L}[\mathbf{Q}]$

$$Q_{(s, a)} \leftarrow Q_{(s, a)} - \alpha \left(r_t + \gamma Q_{(s_{t+1}, a^*)} - Q_{(s_t, a_t)} \right) \left(\gamma \delta(s \stackrel{a}{=} s_{t+1}) - \delta(s \stackrel{a}{=} s_t) \right)$$

$$Q_{(s_t, a_t)} \leftarrow (1 - \alpha) Q_{(s_t, a_t)} + \alpha \left(r_t + \gamma Q_{(s_{t+1}, a^*)} \right) \quad \checkmark$$

$$Q_{(s_{t+1}, a^*)} \leftarrow (1 - \alpha \gamma^2) Q_{(s_{t+1}, a^*)} - \alpha \gamma \left(r_t - Q_{(s_t, a_t)} \right) \quad \times$$

- $Q_{(s_{t+1}, a^*)}$ does not depend on $Q_{(s_t, a_t)}$!
 - residual gradients break causality
 - in practice slows down learning considerably

residual algorithms converge, but not necessarily to the correct values, as s_t and s_{t+1} are not independent (Baird, 1995)



core concept: Bootstrapping

- causality: future values do not depend on past values
- separate bootstrapping network $v_{\theta'}$ with parameters θ'

$$\mathcal{L}[\theta] \quad := \quad \mathbb{E}_{\mathcal{D}} \left[\left(\underbrace{r_t + \gamma v_{\theta'}(s_{t+1})}_{\text{bootstrapped target}} - v_{\theta}(s_t) \right)^2 \right]$$



core concept: Bootstrapping

- causality: future values do not depend on past values
- separate bootstrapping network $v_{\theta'}$ with parameters θ'

$$\mathcal{L}[\theta] := \mathbb{E}_{\mathcal{D}} \left[\left(\underbrace{r_t + \gamma v_{\theta'}(s_{t+1})}_{\text{bootstrapped target}} - v_{\theta}(s_t) \right)^2 \right]$$

- Semi-gradient TD-learning: $\theta' = \theta$
 - `targets.detach()` prevents gradient propagation
 - learns much faster
 - convergence for linear models

[Gordon \(1995\)](#) showed faster learning; for the convergence proof see [Tsitsiklis and Van Roy \(1997\)](#)



core concept: Bootstrapping

- causality: future values do not depend on past values
- separate bootstrapping network $v_{\theta'}$ with parameters θ'

$$\mathcal{L}[\theta] := \mathbb{E}_{\mathcal{D}} \left[\left(\underbrace{r_t + \gamma v_{\theta'}(s_{t+1})}_{\text{bootstrapped target}} - v_{\theta}(s_t) \right)^2 \right]$$

- Semi-gradient TD-learning: $\theta' = \theta$
 - `targets.detach()` prevents gradient propagation
 - learns much faster
 - convergence for linear models
- Neural-fitted Q-iteration (NFQ): update $\theta' \leftarrow \theta$ after convergence
 - first successful deep RL algorithm, but slow iterations

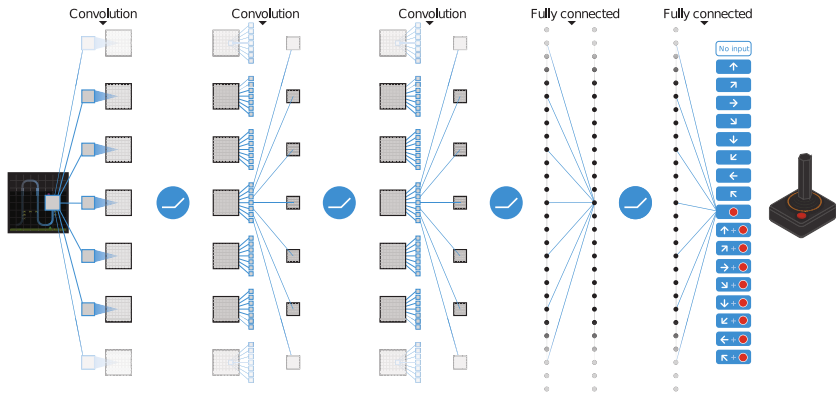
Neural-fitted Q-iteration by [Riedmiller \(2005\)](#)

3.1 Online Q-learning



- Semi-gradient Q-learning using NN with one output per action

$$q_{\theta}(s_t, a_t) := \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t, a_t \right] \equiv r(s_t, a_t) + \gamma \mathbb{E} \left[\max_{a'} q_{\theta}(s_{t+1}, a') \right]$$



originally by [Watkins and Dayan \(1992\)](#), deep version by [Mnih et al. \(2013, figure source\)](#), see [Sutton and Barto \(2018\)](#) for details

3.1 Online Q-learning



- Semi-gradient Q-learning using NN with one output per action

$$\underbrace{q_{\theta}(s_t, a_t)}_{\text{value}} := \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \middle| \begin{smallmatrix} s_t \\ a_t \end{smallmatrix} \right] \equiv \underbrace{r(s_t, a_t) + \gamma \mathbb{E} \left[\max_{a'} q_{\theta}(s_{t+1}, a') \right]}_{\text{target}}$$

End-to-end regression of Q-value function q in environment

```
1 from torch.optim import RMSprop
2 from torch.nn.functional import mse_loss
3 optimizer = RMSprop(q.parameters())
4 while True: # sample episode from environment
5     state, action, reward, term, next = environment.sample(q)
6     # compute left and right side of Bellman eq.
7     value = q(state).gather(dim=-1, index=action)
8     target = reward + gamma * (~term * q(next).max(dim=-1)[0])
9     # gradient descent step on supervised regression loss
10    optimizer.zero_grad()
11    mse_loss(value, target.detach()).backward()
12    optimizer.step()
```

originally by [Watkins and Dayan \(1992\)](#), deep version by [Mnih et al. \(2013\)](#), figure source), see [Sutton and Barto \(2018\)](#) for details

3.1 Online Q-learning



- Semi-gradient Q-learning using NN with one output per action

$$\underbrace{q_{\theta}(s_t, a_t)}_{\text{value}} := \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \middle| \begin{smallmatrix} s_t \\ a_t \end{smallmatrix} \right] \equiv \underbrace{r(s_t, a_t) + \gamma \mathbb{E} \left[\max_{a'} q_{\theta}(s_{t+1}, a') \right]}_{\text{target}}$$

End-to-end regression of Q-value function q in environment

```
1 from torch.optim import RMSprop
2 from torch.nn.functional import mse_loss
3 optimizer = RMSprop(q.parameters())
4 while True: # sample episode from environment
5     state, action, reward, term, next = environment.sample(q)
6     # compute left and right side of Bellman eq.
7     value = q(state).gather(dim=-1, index=action)
8     target = reward + gamma * (~term * q(next).max(dim=-1)[0])
9     # gradient descent step on supervised regression loss
10    optimizer.zero_grad()
11    mse_loss(value, target.detach()).backward()
12    optimizer.step()
```

- term** boolean indicates terminal state. Why do we need this?

originally by [Watkins and Dayan \(1992\)](#), deep version by [Mnih et al. \(2013\)](#), figure source), see [Sutton and Barto \(2018\)](#) for details

- Residual value gradients are slow
- We use in practice semi-gradients with `.detach()`
- Bootstrapping varies from semi-gradients to NFQ
- Q-networks have one head per action

Learning Objectives

LO3.1: Derive and implement semi-gradients of a (Q-)value function

LO3.2: Define and implement Q-network architectures

3.2

Deep Q-Learning Stabilization Techniques

3.2 Catastrophic forgetting



- Online learning violates ML assumptions
 - regression targets (bootstrapping) not stationary
 - greedy policy changes training distribution
 - networks “forget” old samples
 - transitions are not i.i.d.




modified from image source: [wikipedia.org](https://en.wikipedia.org/wiki/Catastrophic_forgetting)

3.2 Catastrophic forgetting



- Online learning violates ML assumptions
 - regression targets (bootstrapping) not stationary
 - greedy policy changes training distribution
 - networks “forget” old samples
 - transitions are not i.i.d.
- Naive online Q-learning is not stable!



experiment on Cartpole-v0



modified from image source: [wikipedia.org](https://en.wikipedia.org/wiki/Catastrophic_forgetting)

3.2 On- and Off-policy sampling



- No training/test set split in RL
 - all data is sampled “live” from environment
 - instead: “who” sampled episodes?

3.2 On- and Off-policy sampling



- No training/test set split in RL
 - all data is sampled “live” from environment
 - instead: “who” sampled episodes?
- On-policy sampling evaluates the sampled policy π
 - + relatively stable and low-dimensional
 - must resample once policy changes

$$V^\pi(s_t) = \mathbb{E} \left[r_t + \gamma V^\pi(s_{t+1}) \mid \overset{a \sim \pi(\cdot | s_t)}{r_t = r(s_t, a)}, s_{t+1} \sim P(\cdot | s_t, a) \right]$$

3.2 On- and Off-policy sampling



- No training/test set split in RL
 - all data is sampled “live” from environment
 - instead: “who” sampled episodes?
- On-policy sampling evaluates the sampled policy π
 - + relatively stable and low-dimensional
 - must resample once policy changes

$$V^\pi(s_t) = \mathbb{E} \left[r_t + \gamma V^\pi(s_{t+1}) \mid \begin{array}{c} a \sim \pi(\cdot | s_t) \\ r_t = r(s_t, a), s_{t+1} \sim P(\cdot | s_t, a) \end{array} \right]$$



core concept: Off-policy sampling evaluates *any* policy π

- larger input space (actions)
- + can reuse old/other's experiences

$$Q^\pi(s_t, a_t) = \mathbb{E} \left[r_t + \gamma Q^\pi(s_{t+1}, a') \mid \begin{array}{c} r_t = r(s_t, a_t), s_{t+1} \sim P(\cdot | s_t, a_t) \\ a' \sim \pi(\cdot | s_{t+1}) \end{array} \right]$$

3.2 Experience replay



- How can we exploit that Q-learning is **off-policy**?



image source: [wikipedia.org](https://www.wikipedia.org)

3.2 Experience replay



- How can we exploit that Q-learning is **off-policy**?



core concept: Experience replay buffers

- remember the last n transitions
- mini-batches of i.i.d. transitions
- always include last episode



experience replay buffers introduced by [Lin \(1992\)](#)

3.2 Experience replay



- How can we exploit that Q-learning is **off-policy**?



core concept: Experience replay buffers

- remember the last n transitions
 - mini-batches of i.i.d. transitions
 - always include last episode
-
- Prioritized experience replay buffers
 - choose transitions with high errors more often
 - more susceptible to catastrophic forgetting




basic idea called *prioritized sweeping* ([Moore and Atkeson, 1993](#)), prioritized replay introduced by ([Schaul et al., 2015](#))

3.2 Exploration and neural networks



- Recap: ϵ -greedy exploration policy
 - act randomly with probability ϵ
 - act greedily with probability $1 - \epsilon$
 - linearly decay ϵ over n_ϵ steps




more on exploration in Lecture 5, image source: www.wikipedia.org

3.2 Exploration and neural networks



- Recap: ϵ -greedy exploration policy
 - act randomly with probability ϵ
 - act greedily with probability $1 - \epsilon$
 - linearly decay ϵ over n_ϵ steps
- Greedy policy \rightarrow catastrophic forgetting
 - buffer fills with the same state-actions
 - other actions “forget” their value
 - maximum can select “drifting” actions




more on exploration in Lecture 5, image source: www.wikipedia.org

3.2 Exploration and neural networks



- Recap: ϵ -greedy exploration policy
 - act randomly with probability ϵ
 - act greedily with probability $1 - \epsilon$
 - linearly decay ϵ over n_ϵ steps
- Greedy policy \rightarrow catastrophic forgetting
 - buffer fills with the same state-actions
 - other actions “forget” their value
 - maximum can select “drifting” actions
- Never stop exploring!
 - stop decay at some ϵ_{\min}
 - measure greedy policy in test episodes





more on exploration in Lecture 5, image source: www.wikipedia.org

3.2 Target networks





- Semi-gradient Q-learning
 - with replay buffer
 - update changes targets
 - unstable learning
 - example `Cartpole-v0`



3.2 Target networks



- Semi-gradient Q-learning
 - with replay buffer
 - update changes targets
 - unstable learning
 - example `Cartpole-v0`
- Hard target update
 - use “old” values as targets
 - every n steps: $\theta' \leftarrow \theta$
 - here $n = 10$





assignment sheet 2

hard target updates first in [Mnih et al. \(2013\)](#)

3.2 Target networks



- Semi-gradient Q-learning
 - with replay buffer
 - update changes targets
 - unstable learning
 - example `Cartpole-v0`
- Hard target update
 - use “old” values as targets
 - every n steps: $\theta' \leftarrow \theta$
 - here $n = 10$
- Soft target update
 - $\theta' \leftarrow (1 - \eta)\theta' + \eta\theta$
 - here $\eta = 0.1$



assignment sheet 2

soft target updates first in [Lillicrap et al. \(2016\)](#), hard target updates first in [Mnih et al. \(2013\)](#)

- Values can only be estimated on-policy, Q-values off-policy
- Naive online deep Q-learning is *not stable*
- Catastrophic forgetting requires regular visitations
- Stabilizes with replay buffers and target networks

Learning Objectives

LO3.3: Explain the difference between on-policy and off-policy sampling

LO3.4: Explain why experience replay buffers and target networks stabilize

3.3

Deep Q-Learning

Deep Q-Networks

core concept: Deep Q-networks (DQN)

- Denotes a *family* of online deep Q-learning algorithms
 - originally hard target update, 1 update/step, visual input
- Samples either an episode or n steps between m updates


```
1 # Using mini-'batch' of transitions from replay buffer
2 batch = self.replay_buffer.sample()
3 targets = batch['rewards'] + self.gamma * (~batch['terminals'] \
4         * self.target_q(batch['next_states']).max(dim=-1)[0])
5 values = q(batch['states']).gather(dim=-1, index=batch['actions'])
6 # Backpropagate loss
7 self.optimizer.zero_grad()
8 mse_loss(values, targets.detach()).backward()
9 self.optimizer.step()
10 # Update target network (hard or soft)
11 self.target_model_update()
```



3.3 Update frequency




- Sample efficiency?
 - example `Cartpole-v0`
 - constant exploration $\epsilon = 0.1$



1 Sampling entire episodes

- 1 update/episode




mean and standard deviation over 10 histogram-smoothed seeds

3.3 Update frequency




- Sample efficiency?
 - example `Cartpole-v0`
 - constant exploration $\epsilon = 0.1$



1 Sampling entire episodes

- n updates/episode
- faster learning for $n \leq 256$
- less stable for $n > 16$




mean and standard deviation over 10 histogram-smoothed seeds

3.3 Update frequency



- Sample efficiency?
 - example `Cartpole-v0`
 - constant exploration $\epsilon = 0.1$




1 Sampling entire episodes

- n updates/episode
- faster learning for $n \leq 256$
- less stable for $n > 16$

2 Sampling n steps/update

- old episodes continued
- faster learning
- less stable for $n < 25$

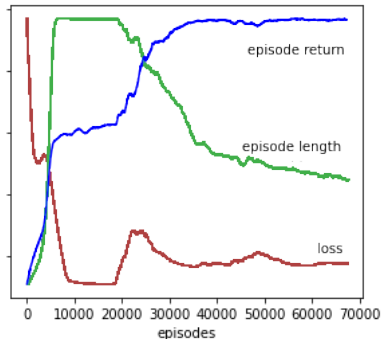


mean and standard deviation over 10 histogram-smoothed seeds

3.3 DQN example: Lunar Lander



- Lunar lander on the moon
 - state: pos., vel. and rot.
 - actions exhaust costly fuel
 - reward for landing between flags
 - punishment for crashing
- 3 phases of learning
 - 1 random initialization
 - 2 learned to hover
 - 3 fine-tune landing
- Loss often increases
 - whenever sampling changes
 - e.g. phase transitions



assignment sheet 2

- **failure:** policy does not pick up
 - raise learning rate α of gradient descent
 - raise γ to consider far future rewards
 - transform network inputs to zero-mean and unit variance
 - transform reward to be in $[-1, 1] \subset \mathbb{R}$ (be careful!)
 - increase exploration time
 - increase episode length (if possible)
 - decrease **or** increase number of layers/neurons
- **instability:** policy unlearns after a while
 - lower learning rate α of gradient descent
 - lower γ to reduce error propagation
 - slow target network adaptation
 - increase mini-batch size
 - increase replay buffer size
 - increase final exploration ϵ

- DQN stabilizes with replay buffers and target networks
- Update frequency crucial for efficiency/stability
- Learning happens in phases, indicated by increasing losses
- Choosing working hyper-parameters requires a lot of experience

Learning Objectives

LO3.5: Implement and evaluate DQN with replay buffer and target network

- Next lecture: **why do we call it *deep learning*?**
- This Thursday is **tutorial**
 - submit **assignment 1!**
 - until start of tutorial
 - answer can be incorrect
- Questions? Ask them here:
answers.ewi.tudelft.nl



image source: xkcd.com



- Leemon Baird. Residual algorithms: Reinforcement learning with function approximation. *Machine Learning*, 1995.
- Geoffrey J. Gordon. Stable function approximation in dynamic programming. In *Proc. 12th International Conference on Machine Learning*, pages 261–268, 1995.
- Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2016. URL <http://arxiv.org/abs/1509.02971>.
- Long-Ji Lin. Self-improving reactive agents based on reinforcement learning, planning and teaching. *Machine Learning*, 8(3): 293–321, 1992.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing atari with deep reinforcement learning. *CoRR*, abs/1312.5602, 2013. URL <http://arxiv.org/abs/1312.5602>. NIPS Deep Learning Workshop 2013.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, Stig Petersen, Charles Beattie, Amir Sadik, Ioannis Antonoglou, Helen King, Dharmashan Kumaran, Daan Wierstra, Shane Legg, and Demis Hassabis. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533, February 2015.
- A.W. Moore and C.G. Atkeson. Prioritized sweeping: Reinforcement learning with less data and less time. *Machine Learning*, 13:103–130, 1993.
- Martin Riedmiller. Neural fitted q iteration – first experiences with a data efficient neural reinforcement learning method. In *Machine Learning: ECML 2005*, pages 317–328. Springer Berlin Heidelberg, 2005. ISBN 978-3-540-31692-3.
- Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized experience replay. In *International Conference on Learning Representations (ICLR)*, 2015. URL <http://arxiv.org/abs/1511.05952>.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. The MIT Press, second edition, 2018. URL <http://incompleteideas.net/book/the-book-2nd.html>.
- John N. Tsitsiklis and Benjamin Van Roy. Analysis of temporal-difference learning with function approximation. In *Advances in Neural Information Processing Systems*, 1997.
- Christopher Watkins and Peter Dayan. Q-learning. *Machine Learning*, 8:279–292, 1992.